Review of P13 — MS1

Summary

This project predicts the probability that one pair beats the other for every scheduled main draw match on the Premier Padel and FIP tour. Additionally, it shows which pair is favoured and by how much.

Strengths

- The data comes from one trusted source, padelapi.org.
- You have a strong baseline. The model has to beat the provider's own Elo, reported to be 71.8% accuracy, 0.517 log loss. The proposal says "if the model only matches the provider's rating, it has added nothing." Some projects pick a weak baseline.
- The provider's Elo is registered first as the champion, therefore predictions exist from the start. A trained model will only replace it if its log loss on the same test set is lower. This makes the promotion rule more explicit.
- From my understanding Elo rates individual players, not pairs. You want to test if the model gets better predictions by adding information like partnership length, form and rest. The project asks a testable question: do partnership length, form and rest add information beyond an Elo rating?
- Assigning pair A and B by a hash of the pair names is a good choice. For example if pair A were always the high seeded pair, then A would win most of the time and then the model would score well by always guessing A.

Concerns or gaps

- The test set might be too small. 20% of 10′000 matches (from Section 3) is roughly 2000 matches (my estimate). "Log loss lower than the baseline" could be achieved by a tiny difference, which could be just noise. The success criterion has no minimum margin.
- The baseline comes from the provider's own 7500 matches, not from your test set, so the comparison may not be fair.

One concrete suggestion

The success criterion: "model log loss lower than the baseline's", but the baseline number comes from a different set of matches. Before MS2, compute the provider's Elo log loss on exactly the same test matches as your model.